


Balanced Search: The Problem


Why are trees important?
Insertion/deletion fast (on every operation, unlike hash table, which has to expand from time to time).
Range queries, sorting (unlike hash tables)
Performance from binary search tree requires remaining balanced (balanced by some constant > 1 at each node).
Ideally, that tree be "bushy"
Leaves (most inner nodes with one child) perform like linked lists
Difference in heights of any two subtrees of a node always differ by at most constant factor K .

Example of Direct Approach: B-Trees



grows/shrinks only at root, then two sides always have
Except root, has from $\lceil M/2 \rceil$ to M children, and one key
Each two children.
From 2 to M children (in non-empty tree).
Added just above bottom: split overfull nodes as needed,
Merge up to parent.


Inserting in B-tree (Simple Case)



CS61B Lecture #29


Search structures (DS(IJ), Chapter 9)
B-trees (DS(IJ), Chapter 11)

Example of Direct Approach: B-Trees



A B-tree is an M -ary search tree, $M > 2$.
B-tree property:
Keys are sorted in each node.
All subtrees to left of a key, K , are $< K$, and all to right
are $> K$.
Leaves at bottom of tree are all empty (don't really exist) and
are all at same level from root.
B-tree is a simple generalization of binary search.

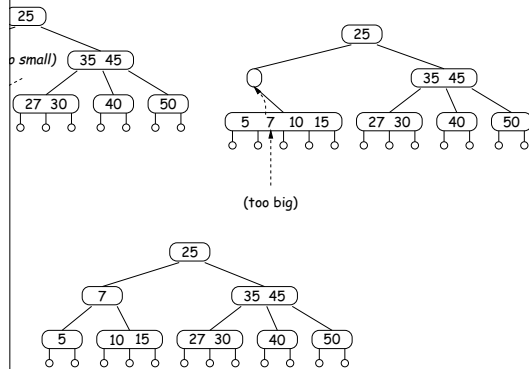
Example Order 4 B-tree ((2,4) Tree)



How to show path when finding 40.
On left side of each child pointer in path bracket 40.
Each node has at least 2 children, and all leaves (little circles) are
at same level, so height must be $O(\lg N)$.
B-tree, order typically much bigger
Depends on size of disk sector, page, or other convenient unit

Deleting Keys from B-tree


from last tree.



13-04 2018

CS61B: Lecture #29 8

Red-Black Tree Constraints



(conceptually) colored red or black.


- Node contains no data (as for B-trees) and is black.
- Every leaf has same number of black ancestors.
- Internal node has two children.
- Red node has two black children.
- 5, and 6 guarantee $O(\lg N)$ searches.

13-04 2018

CS61B: Lecture #29 10

Constraints: Left-Leaning Red-Black Trees

(2,4) or (2,3) tree with three children may be represented in different ways in a red-black tree:



considerably simplify insertion and deletion in a red-black tree by choosing the option on the left.

constraint, there is a one-to-one relationship between (2,4) trees and red-black trees.

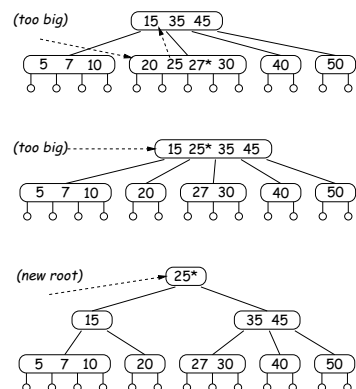
Left-leaning trees are called *left-leaning red-black trees*.

For simplification, let's restrict ourselves to red-black trees that correspond to (2,3) trees (whose nodes have no more than two children), so that no red-black node has two red children.

13-04 2018

CS61B: Lecture #29 12

Inserting in B-Tree (Splitting)



13-04 2018

CS61B: Lecture #29 7

Red-Black Trees

A red-black tree is a binary search tree with additional constraints that prevent it from being too unbalanced if it can be.

Searching is always $O(\lg N)$.

Java's TreeSet and TreeMap types.

When keys are inserted or deleted, tree is *rotated* and *recolor*ed to restore balance.

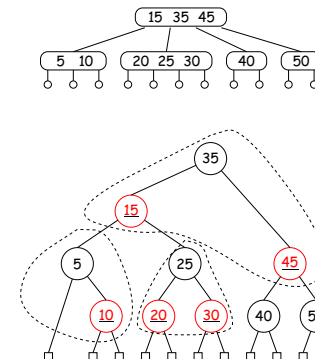
13-04 2018

CS61B: Lecture #29 9

Red-Black Trees and (2,4) Trees

A (2,4) tree corresponds to a (2,4) tree, and the operations on (2,4) trees correspond to those on the other.

Each (2,4) tree corresponds to a cluster of 1-3 red-black trees, with the top node is black and any others are red.




13-04 2018

CS61B: Lecture #29 11

Rotations and Recolorings

cases, we'll augment the general rotation algorithms with coloring.

the color from the original root to the new root, and color the new root red. Examples:



these changes the number of black nodes along any path from root and the leaves.

The Algorithm (Sedgwick)

Binary-tree type `RBTree`: basically ordinary BST nodes


the same as for ordinary BSTs, but we add some fixups to maintain the red-black properties.

```
RBTree tree, KeyType key) {
    if (tree == null)
        return new RBTree(key, null, null, RED);
    int cmp = key.compareTo(tree.label());
    if (cmp < 0) tree.setLeft(insert(tree.left(), key));
    else tree.setRight(insert(tree.right(), key));
}
```

`fixup(tree);` // Only line that's all new!

Fixing Up the Tree (II)

break up 4-nodes into 3-nodes or 2-nodes.




as a result of other fixups, or of insertion into the empty tree, it may end up red, so color the root black after the rest of the fixups are finished. (Not part of the fixup function; the end).

Red-Black Insertion and Rotations

from just as for binary tree (color red except when tree is black).

(and recolor) to restore red-black property, and thus maintain balance.


Red-black trees *preserves* binary tree property, but changes balance.



Splitting by Recoloring

nodes will temporarily create nodes with too many children, but we can fix them up.

Recoloring allows us to split nodes. We'll call it `colorFlip`:




`colorFlip` joins the parent node, splitting the original.

Fixing Up the Tree


to break up the BST, we restore the left-leaning red-black property and limit ourselves to red-black trees that correspond to a valid BST by applying the following (in order) to each node:

convert right-leaning trees to left-leaning:




node B will be red, so that both B and D end up red. This is the end of the fixup function.

convert linked red nodes into a normal 4-node (temporarily).




Insertion Example (II)

), let's insert 6, leading to the tree on the left. This is
 , so apply Fixup 1:




Insertion Example (IIIa)

xup 2.




Insertion Example (IIIc)

another 4-node, so apply fixup 3 again.




of Left-Leaning 2-3 Red-Black Insertion

initial tree on left. No fixups needed.




Insertion Example (III)

r inserting 85. We need fixup 1 first.



Insertion Example (IIIb)

a 4-node, so apply fixup 3.



Insertion Example (IIId)

a right-leaning tree, so apply fixup 1.

