

CS61B Lecture #29

Today:

- Balanced search structures (*DS(IJ)*, Chapter 9)


Coming Up:

- Pseudo-random Numbers (*DS(IJ)*, Chapter 11)

Balanced Search: The Problem


- Why are search trees important?
 - Insertion/deletion fast (on every operation, unlike hash tables which has to expand from time to time).
 - Support range queries, sorting (unlike hash tables)
- But $O(\lg N)$ performance from binary search tree requires keys be divided \approx by some constant > 1 at each node.
- In other words, that tree be "bushy"
- "Stringy" trees (most inner nodes with one child) perform like lists.
- Suffices that heights of any two subtrees of a node always differ by no more than constant factor K .

Example of Direct Approach: B-Trees



- *Order M B-tree* is an M -ary search tree, $M > 2$.
- Obeys search-tree property:
 - Keys are sorted in each node.
 - All keys in subtrees to left of a key, K , are $< K$, and all are $> K$.
- Children at bottom of tree are all empty (don't really exist equidistant from root).
- Searching is simple generalization of binary search.


Example of Direct Approach: B-Trees



Idea: If tree grows/shrinks only at root, then two sides always same height.

- Each node, except root, has from $\lceil M/2 \rceil$ to M children, and "between" each two children.
- Root has from 2 to M children (in non-empty tree).
- Insertion: add just above bottom; split overfull nodes as moving one key up to parent.


Sample Order 4 B-tree ((2,4) Tree)




- Crossed lines show path when finding 40.
- Keys on either side of each child pointer in path bracket 40
- Each node has at least 2 children, and all leaves (little circles at the bottom, so height must be $O(\lg N)$).
- In real-life B-tree, order typically much bigger
 - comparable to size of disk sector, page, or other convenient unit of I/O

Inserting in B-tree (Simple Case)

- Start:




- Insert 7:




Inserting in B-Tree (Splitting)

- Insert 27:


(too big)



(too big)



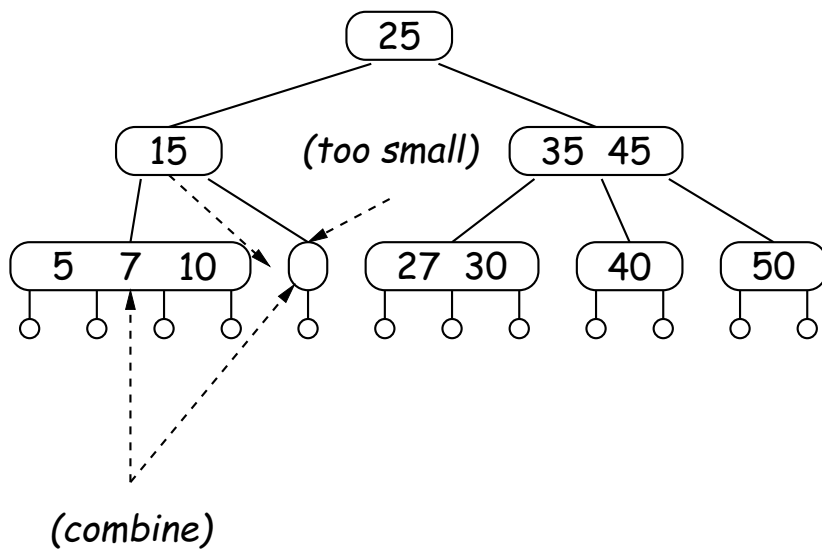


(new root)



Deleting Keys from B-tree


- Remove 20 from last tree.



Red-Black Trees

- Red-black tree is a binary search tree with additional constraints that limit how unbalanced it can be.
- Thus, searching is always $O(\lg N)$.
- Used for Java's TreeSet and TreeMap types.
- When items are inserted or deleted, tree is *rotated* and *rebalanced* as needed to restore balance.



Red-Black Tree Constraints



1. Each node is (conceptually) colored red or black.
 2. Root is black.
 3. Every leaf node contains no data (as for B-trees) and is black.
 4. Every leaf has same number of black ancestors.
 5. Every internal node has two children.
 6. Every red node has two black children.
- Conditions 4, 5, and 6 guarantee $O(\lg N)$ searches.


Red-Black Trees and (2,4) Trees

- Every red-black tree corresponds to a (2,4) tree, and the operations on one correspond to those on the other.
- Each node of (2,4) tree corresponds to a cluster of 1-3 red nodes in which the top node is black and any others are red.



Additional Constraints: Left-Leaning Red-Black T


- A node in a (2,4) or (2,3) tree with three children may be represented in two different ways in a red-black tree:



- We can considerably simplify insertion and deletion in a red-black tree by always choosing the option on the left.
- With this constraint, there is a one-to-one relationship between (2,4) trees and red-black trees.
- The resulting trees are called *left-leaning red-black trees*.
- As a further simplification, let's restrict ourselves to red-black trees that correspond to (2,3) trees (whose nodes have at most 3 children), so that no red-black node has two red children.


Red-Black Insertion and Rotations

- Insert at bottom just as for binary tree (color red except when initially empty).
- Then rotate (and recolor) to restore red-black property, and balance.
- Rotation of trees *preserves* binary tree property, but changes balance.



Rotations and Recolorings





- For our purposes, we'll augment the general rotation algorithm with some recoloring.
- Transfer the color from the original root to the new root, and the original root red. Examples:



- Neither of these changes the number of black nodes along a path between the root and the leaves.

Splitting by Recoloring

- Our algorithms will temporarily create nodes with too many children and then split them up.
- A simple recoloring allows us to split nodes. We'll call it **color**.



- Here, key 10 joins the parent node, splitting the original.


The Algorithm (Sedgewick)

- We posit a binary-tree type `RBTree`: basically ordinary BST plus color.
- Insertion is the same as for ordinary BSTs, but we add some code to restore the red-black properties.

```
RBTree insert(RBTree tree, KeyType key) {  
    if (tree == null)  
        return new RBTree(key, null, null, RED);  
    int cmp = key.compareTo(tree.label());  
    else if (cmp < 0) tree.setLeft(insert(tree.left(  
    else  
        tree.setRight(insert(tree.right(  
  
    return fixup(tree);    // Only line that's all  
}
```


Fixing Up the Tree


- As we return back up the BST, we restore the left-leaning red properties, and limit ourselves to red-black trees that correspond to (2,3) trees by applying the following (in order) to each node:
- Fixup 1: Convert right-leaning trees to left-leaning:



```
if (tree.right().isRed()
    && tree.left().isRed())
    tree.rotateLeft();
}
```

Sometimes, node B will be red, so that both B and D end up red. This is fixed by...


- Fixup 2: Rotate linked red nodes into a normal 4-node (temporarily breaking the BST property)



```
if (tree.left().isRed()
    && tree.left().left().isRed())
    tree.rotateRight();
```

Fixing Up the Tree (II)



- Fixup 3: Break up 4-nodes into 3-nodes or 2-nodes.



- Fixup 4: As a result of other fixups, or of insertion into the tree, the root may end up red, so color the root black after - of insertion and fixups are finished. (Not part of the fixup f just done at the end).



Example of Left-Leaning 2-3 Red-Black Insert

- Insert 0 into initial tree on left. No fixups needed.





Insertion Example (II)

- Instead of 0, let's insert 6, leading to the tree on the left. right-leaning, so apply Fixup 1:





Insertion Example (III)

- Now consider inserting 85. We need fixup 1 first.





Insertion Example (IIIa)

- Now apply fixup 2.





Insertion Example (IIIb)

- This gives us a 4-node, so apply fixup 3.



Insertion Example (IIIc)

- This gives us another 4-node, so apply fixup 3 again.



Insertion Example (IIId)

- This gives us a right-leaning tree, so apply fixup 1.

