

Chapter 9

Coming Up:

- Pseudo-random Numbers (*DS(IJ)*, Chapter 11)

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 1

- Haven't said much about cost of comparisons.
- For strings, worst case is length of string.
- Therefore should throw extra factor of key length, L , into costs:
 - $\Theta(M)$ comparisons really means $\Theta(ML)$ operations.
 - So to look for key X , keep looking at same chars of X M times.
- Can we do better? Can we get search cost to be $O(L)$?

Idea: Make a *multi-way decision tree*, with one decision per character of key.

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 2

{a, abase, abash, abate, abbas, axolotl, axe, fabric, facet}

- Ticked lines show paths followed for "abash" and "fabric"
- Each internal node corresponds to a possible prefix.
- Characters in path to node = that prefix.

Last modified: Thu Nov 1 19:39:39 2018

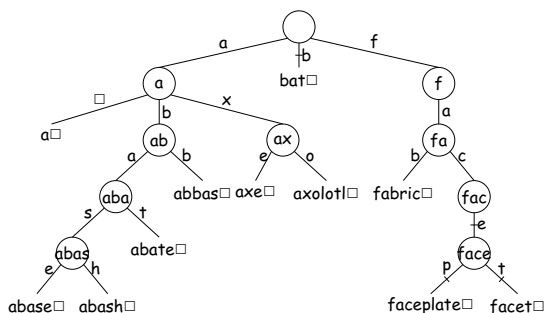
CS61B: Lecture #31 3



Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 4

- New edges ticked.



Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 5

Internal nodes is array indexed by character.

- Gives $O(L)$ performance, L length of search key.
- [Looks as if independent of N , number of keys. Is there a dependence?]
- Problem:** arrays are *sparsely populated* by non-null values—waste of space.

Idea: Put the arrays on top of each other!


- Use null (0, empty) entries of one array to hold non-null elements of another.
- Use extra markers to tell which entries belong to which array.

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 6

ceasing slides)

- Three leaf arrays, each indexed 0..9



Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 7

- Not so good if we want to expand our trie.
- A bit complicated.
- Actually more useful for representing large, sparse, fixed tables with many rows and columns.
- Furthermore, number of children in trie tends to drop drastically when one gets a few levels down from the root.
- So in practice, might as well use linked lists to represent set of node's children...
- ...but use arrays for the first few levels, which are likely to have more children.

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 8

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 9

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 10

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 11

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 12

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 13

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 14

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 15

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 16

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 17

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 18

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 19

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 20

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 21

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 22

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 23

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 24

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 25

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 26

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:



- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 27

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 28

n-ary search tree in which we choose to put the keys at "random" heights.

- More often thought of as an ordered list in which one can skip large segments.
- Typical example:




- To search, start at top layer on left, search until next step would overshoot, then go down one layer and repeat.
- In list above, we search for 125 and 127. Gray nodes are looked at; darker gray nodes are overshoots.
- Heights of the nodes were chosen randomly so that there are about 1/2 as many nodes

Last modified: Thu Nov 1 19:39:39 2018


CS61B: Lecture #31 29

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 30



- In any order, we add 126 and 127 (choosing random heights for them), and remove 20 and 40:



- Shaded nodes here have been modified.

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 31

$\Theta(\lg N)$ performance.

- B-trees, red-black trees:
 - Give $\Theta(\lg N)$ performance for searches, insertions, deletions.
 - B-trees good for external storage. Large nodes minimize # of I/O operations
- Tries:
 - Give $\Theta(B)$ performance for searches, insertions, and deletions, where B is length of key being processed.
 - But hard to manage space efficiently.
- Interesting idea: scrunched arrays share space.
- Skip lists:
 - Give probable $\Theta(\lg N)$ performance for searches, insertions, deletions

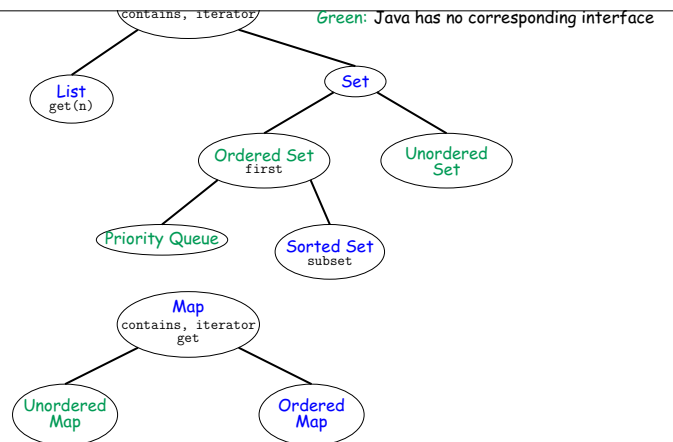
Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 32

tures.

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 33



Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 34

Multiset

- List: arrays, linked lists, circular buffers
- Set
 - OrderedSet
 - * Priority Queue: heaps
 - * Sorted Set: binary search trees, red-black trees, B-trees, sorted arrays or linked lists
 - Unordered Set: hash table

Map

- Unordered Map: hash table
- Ordered Map: red-black trees, B-trees, sorted arrays or linked lists

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 35

- List: ArrayList, LinkedList, Stack, ArrayBlockingQueue, ArrayDeque
- Set
 - OrderedSet
 - * Priority Queue: PriorityQueue
 - * Sorted Set (SortedSet): TreeSet
 - Unordered Set: HashSet

Map

- Unordered Map: HashMap
- Ordered Map (SortedMap): TreeMap

Last modified: Thu Nov 1 19:39:39 2018

CS61B: Lecture #31 36